Our module

- Search & Recommendations APIs
- Storefront owns customer UI/UX
- Exercise task: name search with category and price filters

A context-package exercise, not a complete shop implementation

A real interface mismatch

- Admin: article_id, category_id, tags, is_published
- Marketplace: productId, category, attributes, status
- Ownership and shared IDs are still being discussed
- Our revised context records both proposals

Joint review is open; no claim that both teams have finalised their files

Two fresh sessions

- Same model and reasoning: gpt-6-astra, medium
- Same task, starter files and six fictional products
- A: task only B: task + original 663-token context
- Original evidence preserved; 18 independent checks

Current module: 690 tokens; the historical experiment was not rerun on it

Both passed; conventions differed

- Independent checks: A 18/18 | B 18/18
- Agent-written tests: A 14 passed | B 13 passed
- Both: 3 messages, 5 commands, 1 file-change batch
- A: query / products / manual validation
- B: q / items / Zod / kebab-case filenames

Zero human code corrections; A was not given the module conventions

What we learned

- Context improved convention alignment in this pair
- Both runs met the common functional requirements
- No observed reduction in messages or commands
- Keep context current and preserve experiment versions

Next: settle shared IDs/ownership and confirm both teams update their files